

# BUSINESS ANALYTICS ASSIGNMENT\_2

ALLEN RICHARDS

2023-10-15

```
library(dplyr)
```

```
##  
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':  
##  
##   filter, lag
```

```
## The following objects are masked from 'package:base':  
##  
##   intersect, setdiff, setequal, union
```

```
library(magrittr)  
library(zoo)
```

```
##  
## Attaching package: 'zoo'
```

```
## The following objects are masked from 'package:base':  
##  
##   as.Date, as.Date.numeric
```

```
library(readr)
```

```
# Load the dataset  
retail_data <- read.csv("D:/KSU SEM-1/BUSINESS ANALYTICS/Assignment_2/OnlineRetail.csv", head  
er = TRUE, stringsAsFactors = FALSE)
```

Question 1 - Show the breakdown of the number of transactions by countries i.e., how many transactions are in the dataset for each country (consider all records including cancelled transactions). Show this in total

number and also in percentage. Show only countries accounting for more than 1% of the total transactions

```
OnlineRetail <- read.csv("D:/KSU SEM-1/BUSINESS ANALYTICS/Assignment_2/OnlineRetail.csv")

Countries_count = OnlineRetail %>% group_by(Country) %>% count(Country)
Countries_pct = OnlineRetail %>% group_by(Country) %>% summarise(percent = 100* n()/nrow(OnlineRetail))
Fltrd_Cntry_pct = filter(Countries_pct, percent>1)
#Countries Count
Countries_count
```

```
## # A tibble: 38 × 2
## # Groups:   Country [38]
##   Country      n
##   <chr>      <int>
## 1 Australia    1259
## 2 Austria      401
## 3 Bahrain      19
## 4 Belgium     2069
## 5 Brazil       32
## 6 Canada      151
## 7 Channel Islands 758
## 8 Cyprus      622
## 9 Czech Republic 30
## 10 Denmark    389
## # i 28 more rows
```

```
#Percentage of transactions greater than 1
Fltrd_Cntry_pct
```

```
## # A tibble: 4 × 2
##   Country      percent
##   <chr>      <dbl>
## 1 EIRE        1.51
## 2 France      1.58
## 3 Germany     1.75
## 4 United Kingdom 91.4
```

Question 2 Create a new variable 'TransactionValue' that is the product of the existing 'Quantity' and 'UnitPrice' variables. Add this variable to the dataframe.

```
# Create the new variable 'TransactionValue' by multiplying 'Quantity' and 'UnitPrice'
OnlineRetail$TransactionValue <- OnlineRetail$Quantity * OnlineRetail$UnitPrice

# Display the first few rows of the dataframe to verify the new variable
head(OnlineRetail)
```

```
## InvoiceNo StockCode Description Quantity
## 1 536365 85123A WHITE HANGING HEART T-LIGHT HOLDER 6
## 2 536365 71053 WHITE METAL LANTERN 6
## 3 536365 84406B CREAM CUPID HEARTS COAT HANGER 8
## 4 536365 84029G KNITTED UNION FLAG HOT WATER BOTTLE 6
## 5 536365 84029E RED WOOLLY HOTTIE WHITE HEART. 6
## 6 536365 22752 SET 7 BABUSHKA NESTING BOXES 2
## InvoiceDate UnitPrice CustomerID Country TransactionValue
## 1 12/1/2010 8:26 2.55 17850 United Kingdom 15.30
## 2 12/1/2010 8:26 3.39 17850 United Kingdom 20.34
## 3 12/1/2010 8:26 2.75 17850 United Kingdom 22.00
## 4 12/1/2010 8:26 3.39 17850 United Kingdom 20.34
## 5 12/1/2010 8:26 3.39 17850 United Kingdom 20.34
## 6 12/1/2010 8:26 7.65 17850 United Kingdom 15.30
```

Question 3 Using the newly created variable, TransactionValue, show the breakdown of transaction values by countries i.e. how much money in total has been spent each country. Show this in total sum of transaction values. Show only countries with total transaction exceeding 130,000 British Pound

```
# Load the required Libraries
library(dplyr)

# Breakdown of transaction values by countries
Trans_sum = OnlineRetail %>% group_by(Country) %>% summarise(sum=sum(TransactionValue))
Fltrd_Trans_sum = filter(Trans_sum, Trans_sum$sum > 130000)
# Sum of TransactionValue for each countries
Trans_sum
```

```
## # A tibble: 38 × 2
##   Country      sum
##   <chr>      <dbl>
## 1 Australia  137077.
## 2 Austria    10154.
## 3 Bahrain     548.
## 4 Belgium   40911.
## 5 Brazil     1144.
## 6 Canada     3666.
## 7 Channel Islands 20086.
## 8 Cyprus    12946.
## 9 Czech Republic  708.
## 10 Denmark   18768.
## # i 28 more rows
```

```
# Group and summarize the data by country, calculating the total transaction value
country_transaction_values <- OnlineRetail %>%
  group_by(Country) %>%
  summarise(Total_TransactionValue = sum(TransactionValue))

# Filter countries with a total transaction value exceeding £130,000
country_transaction_values_filtered <- country_transaction_values %>%
  filter(Total_TransactionValue > 130000)

# Print the result
print(country_transaction_values_filtered)
```

```
## # A tibble: 6 × 2
##   Country      Total_TransactionValue
##   <chr>      <dbl>
## 1 Australia  137077.
## 2 EIRE      263277.
## 3 France    197404.
## 4 Germany   221698.
## 5 Netherlands 284662.
## 6 United Kingdom 8187806.
```

Question 4 we are dealing with the InvoiceDate variable. The variable is read as a categorical when you read data from the file. Now we need to explicitly instruct R to interpret this as a Date variable. “POSIXlt” and “POSIXct” are two powerful object classes in R to deal with date and time.

```
# First Let's convert 'InvoiceDate' into a POSIXlt object
Temp <- strptime(OnlineRetail$InvoiceDate, format = '%m/%d/%Y %H:%M', tz = 'GMT')


# Check the First few entries of the 'TEMP' variable
head(Temp)
```

```
## [1] "2010-12-01 08:26:00 GMT" "2010-12-01 08:26:00 GMT"
## [3] "2010-12-01 08:26:00 GMT" "2010-12-01 08:26:00 GMT"
## [5] "2010-12-01 08:26:00 GMT" "2010-12-01 08:26:00 GMT"
```

```
# Let Separate date, day of the week, and hour components
OnlineRetail$New_Invoice_Date <- as.Date(Temp)
OnlineRetail$Invoice_Day_Week <- weekdays(OnlineRetail$New_Invoice_Date)
OnlineRetail$New_Invoice_Hour <- as.numeric(format(Temp, "%H"))
OnlineRetail$New_Invoice_Month <- as.numeric(format(Temp, "%m"))

# The Date objects have a lot of flexible functions. For example knowing two date values, the
object allows you to know the difference between the two dates in terms of the number days.

OnlineRetail$New_Invoice_Date[20000] - OnlineRetail$New_Invoice_Date[10]
```



```
## Time difference of 8 days
```

```
# Let Convert Dates to Days of the Week
OnlineRetail$Invoice_Day_Week <- weekdays(OnlineRetail$New_Invoice_Date)

# Let Extract hour component
OnlineRetail$New_Invoice_Hour <- as.numeric(format(Temp, "%H"))

# Let define the month as a seperate numeric variable
OnlineRetail$New_Invoice_Month <- as.numeric(format(Temp, "%m"))
```

(a) Let show the percentage of transaction  
(by number) by days of the week

```
# Assuming you have the 'Online_Retail' dataset with the 'Invoice_Day_Week' variable
# You can use the table() function to count the number of transactions for each day of the week
day_of_week_counts <- table(OnlineRetail$Invoice_Day_Week)

# Calculate the total number of transactions
total_transactions <- sum(day_of_week_counts)

# Calculate the percentage for each day of the week
percentages <- (day_of_week_counts / total_transactions) * 100

# Create a data frame to display the results
result_df <- data.frame(DayOfWeek = names(day_of_week_counts),
                        Count = as.vector(day_of_week_counts),
                        Percentage = percentages)

# Display the result
print(result_df)
```

```
##   DayOfWeek Count Percentage.Var1 Percentage.Freq
## 1   Friday  82193           Friday      15.16731
## 2   Monday  95111           Monday      17.55110
## 3   Sunday  64375           Sunday      11.87930
## 4 Thursday 103857          Thursday      19.16503
## 5   Tuesday 101808          Tuesday      18.78692
## 6 Wednesday 94565          Wednesday      17.45035
```

## (b) Let Show the percentage of transactions (by transaction volume) by days of the week

```
# Assuming you have the 'OnlineRetail' dataset with the 'Invoice_Day_Week' variable
# You can use the aggregate() function to calculate the total transaction volume (sum of Quantity) for each day of the week

OnlineRetail %>% group_by(Invoice_Day_Week)%>% summarise(Volume.of.transaction=(sum(TransactionValue)))%>% mutate(Volume.of.transaction,'percent'=(Volume.of.transaction*100)/sum(Volume.of.transaction))
```

```
## # A tibble: 6 × 3
##   Invoice_Day_Week Volume.of.transaction percent
##   <chr>           <dbl>      <dbl>
## 1 Friday          1540611.    15.8
## 2 Monday          1588609.    16.3
## 3 Sunday           805679.     8.27
## 4 Thursday        2112519     21.7
## 5 Tuesday         1966183.    20.2
## 6 Wednesday       1734147.    17.8
```

## (c) Let Show the percentage of transactions (by transaction volume) by month of the year

```
OnlineRetail %>% group_by(New_Invoice_Month)%>%
summarise(Volume.By.Month=sum(TransactionValue))%>% mutate(Volume.By.Month, 'Percent'=(Volume.
By.Month*100)/sum(Volume.By.Month))
```

```
## # A tibble: 12 × 3
##   New_Invoice_Month Volume.By.Month Percent
##           <dbl>           <dbl>   <dbl>
## 1               1         560000.    5.74
## 2               2         498063.    5.11
## 3               3         683267.    7.01
## 4               4         493207.    5.06
## 5               5         723334.    7.42
## 6               6         691123.    7.09
## 7               7         681300.    6.99
## 8               8         682681.    7.00
## 9               9        1019688.   10.5
## 10             10        1070705.   11.0
## 11             11        1461756.   15.0
## 12             12        1182625.   12.1
```

## Question : What was the date with the highest number of transactions from Australia?

```
# Assuming you have the 'Online_Retail' dataset with 'InvoiceDate' and 'Country' columns

# Filter the dataset to include only transactions from Australia
australia_data <- OnlineRetail[OnlineRetail$Country == "Australia", ]

# Create a table of counts for each unique InvoiceDate within Australia
date_counts <- table(australia_data$InvoiceDate)

# Find the date with the highest number of transactions
max_date <- names(date_counts[date_counts == max(date_counts)])

# Display the result
print(max_date)
```

```
## [1] "6/15/2011 13:37"
```

## Question E : The company needs to shut down the website for two consecutive hours

for maintenance. What would be the hour of the day to start this so that the distribution is at minimum for the customers? The responsible IT team is available from 7:00 to 20:00 every day.

```
# Assuming you have the 'Online_Retail' dataset with 'New_Invoice_Hour' indicating the hour of transactions

# Filter transactions within the available time frame (7:00 to 20:00)
available_hours_data <- OnlineRetail[OnlineRetail$New_Invoice_Hour >= 7 & OnlineRetail$New_Invoice_Hour <= 20, ]

# Create a table of counts for each hour of the day
hourly_counts <- table(available_hours_data$New_Invoice_Hour)

# Calculate the rolling sum of transaction volumes for two consecutive hours
rolling_sum <- sapply(1:(length(hourly_counts) - 1), function(i) {
  sum(hourly_counts[i:(i + 1)])
})

# Find the starting hour with the minimum rolling sum
optimal_start_hour <- which(rolling_sum == min(rolling_sum))

# Display the result
print(paste("Optimal Start Hour for Maintenance:", optimal_start_hour + 6)) # Add 6 to get the actual hour (e.g., 7:00, 8:00, etc.)
```

```
## [1] "Optimal Start Hour for Maintenance: 19"
```

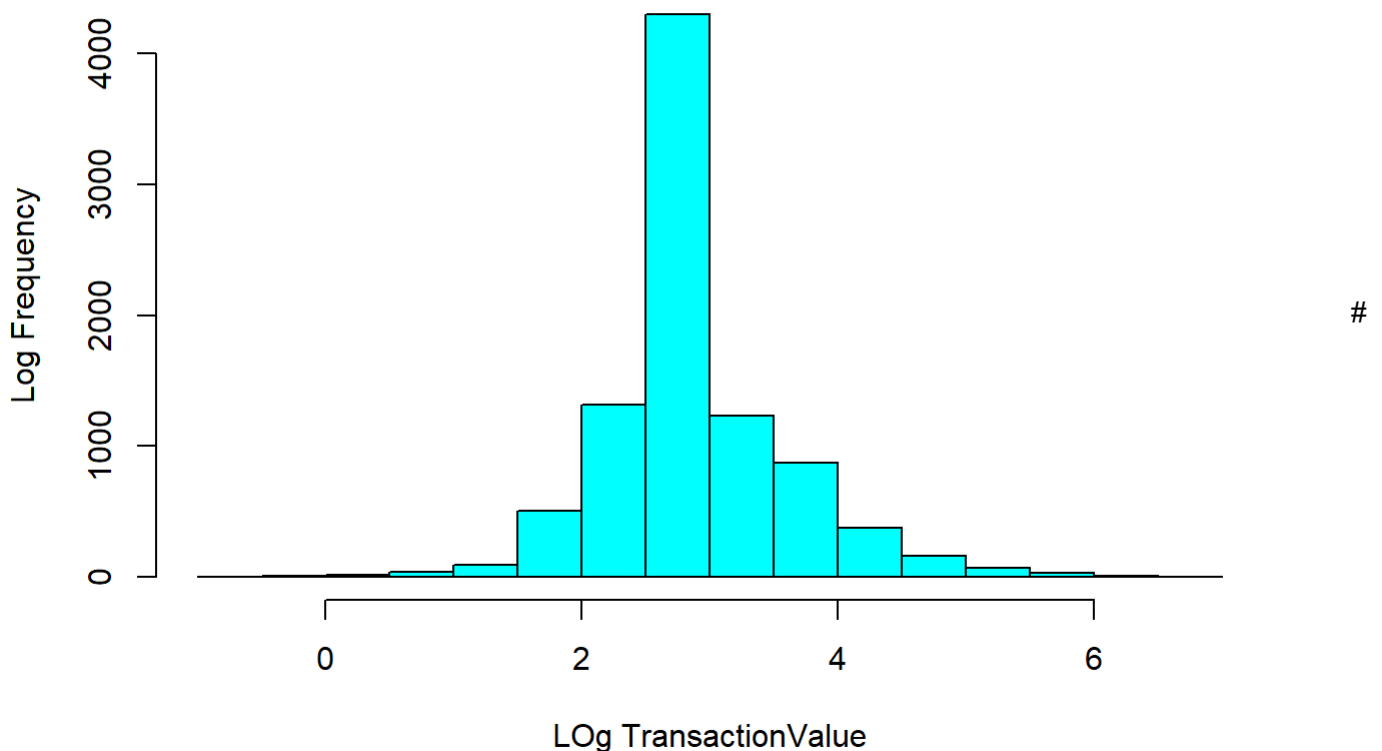
Question 5 : Plot the histogram of transaction values from Germany. Use the hist() function to plot.

```
hist(x=log(OnlineRetail$TransactionValue[OnlineRetail$Country=="Germany"]),xlab = "Log TransactionValue",col = 'cyan' ,main = 'Germany transaction',ylab = 'Log Frequency')
```

```
## Warning in log(OnlineRetail$TransactionValue[OnlineRetail$Country ==
## "Germany"]): NaNs produced
```



## Germany transaction



Question 6 : Which customer had the highest number of transactions? Which customer is most valuable (i.e.highest total sum of transactions)

```
# Let the data summarise the customer transaction
Data123<- OnlineRetail %>% group_by(CustomerID)%>%
summarise(CustomerTransaction = n())%>% filter(CustomerID != "NA")%>% filter(CustomerTransact
ion ==max(CustomerTransaction) )

print(paste('The customerID had the highest number of transactions is',Data123$CustomerID,'wi
th max transaction of ',Data123$CustomerTransaction))
```

```
## [1] "The customerID had the highest number of transactions is 17841 with max transaction o
f 7983"
```

```
Data234<- OnlineRetail %>% group_by(CustomerID)%>%
summarise(total.transaction.by.each.customer = sum(TransactionValue))%>%
arrange(desc(total.transaction.by.each.customer))%>%
filter(CustomerID != "NA")%>% filter(total.transaction.by.each.customer ==max(total.transacti
on.by.each.customer) )

# Calculate the total transaction amount with customer ID
print(paste('Most valuable customerID is',Data234$CustomerID,'with total transaction Amount
$',Data234$total.transaction.by.each.customer))
```

```
## [1] "Most valuable customerID is 14646 with total transaction Amount $ 279489.02"
```

## Question 7 - Calculate the percentage of missing values for each variable in the dataset

```
# Assuming you have the 'OnlineRetail' dataset

# Calculate the percentage of missing values for each variable
missing_percentages <- sapply(OnlineRetail, function(col) {
  missing_count <- sum(is.na(col))
  missing_percentage <- (missing_count / length(col)) * 100
  return(missing_percentage)
})

# Create a data frame to display the missing value percentages
missing_df <- data.frame(Variable = names(missing_percentages), Percentage = missing_percentages)

# Display the result
print(missing_df)
```

##	Variable	Percentage
## InvoiceNo	InvoiceNo	0.00000
## StockCode	StockCode	0.00000
## Description	Description	0.00000
## Quantity	Quantity	0.00000
## InvoiceDate	InvoiceDate	0.00000
## UnitPrice	UnitPrice	0.00000
## CustomerID	CustomerID	24.92669
## Country	Country	0.00000
## TransactionValue	TransactionValue	0.00000
## New_Invoice_Date	New_Invoice_Date	0.00000
## Invoice_Day_Week	Invoice_Day_Week	0.00000
## New_Invoice_Hour	New_Invoice_Hour	0.00000
## New_Invoice_Month	New_Invoice_Month	0.00000

## Question 8 : What are the number of transactions with missing CustomerID records by countries?

```
# Assuming you have the 'Online_Retail' dataset

# Filter the dataset to include only transactions with missing CustomerID
missing_customerID_data <- OnlineRetail[is.na(OnlineRetail$CustomerID), ]

# Create a table of counts for missing CustomerID transactions by country
missing_customerID_counts <- table(missing_customerID_data$Country)

# Display the result
print(missing_customerID_counts)
```

```
##
##      Bahrain      EIRE      France      Hong Kong      Israel
##          2        711         66        288         47
##      Portugal  Switzerland United Kingdom  Unspecified
##          39        125        133600        202
```

## Question 9 : how often the costumers comeback to the website for their next shopping? (i.e. what is the average number of days between consecutive shopping)

```
# Calculate the time difference between consecutive shopping visits for each customer
Average<- OnlineRetail %>% group_by(CustomerID)%>%
summarise(difference.in.consecutivedays= diff(New_Invoice_Date))%>%

# Convert 'CONSECUTIVEDAYS' to a DIFFERENCE
filter(difference.in.consecutivedays>0)
```

```
## Warning: Returning more (or less) than 1 row per `summarise()` group was deprecated in
## dplyr 1.1.0.
## i Please use `reframe()` instead.
## i When switching from `summarise()` to `reframe()`, remember that `reframe()`
## always returns an ungrouped data frame and adjust accordingly.
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was
## generated.
```

```
## `summarise()` has grouped output by 'CustomerID'. You can override using the
## `.groups` argument.
```

```
# Display the result
print(paste('The average number of days between consecutive shopping is',mean(Average$difference.in.consecutivedays)))
```

```
## [1] "The average number of days between consecutive shopping is 38.4875"
```

Question 10 : In the retail sector, it is very important to understand the return rate of the goods purchased by customers. In this example, we can define this quantity, simply, as the ratio of the number of transactions cancelled (regardless of the transaction value) over the total number of transactions. Page 4 With this definition, what is the return rate for the French customers?

```
# Assuming you have the 'Online_Retail' dataset

# Filter the dataset to include only transactions from French customers
french_data <- OnlineRetail[OnlineRetail$Country == "France", ]

# Calculate the total number of transactions for French customers
total_transactions_french <- nrow(french_data)

# Calculate the number of canceled transactions for French customers
canceled_transactions_french <- sum(substr(french_data$InvoiceNo, 1, 1) == "C")

# Calculate the return rate for French customers
return_rate_french <- canceled_transactions_french / total_transactions_french

# Display the result as a percentage
return_rate_percentage <- return_rate_french * 100
print(paste("Return Rate for French Customers:", round(return_rate_percentage, 2), "%"))
```

```
## [1] "Return Rate for French Customers: 1.74 %"
```

Question - 11 What is the product that has generated the highest revenue for the retailer? (i.e. item with the highest total sum of 'TransactionValue').

```
# Assuming you have the 'Online_Retail' dataset with 'StockCode', 'Quantity', and 'UnitPrice' columns

# Calculate the 'TransactionValue' for each transaction
OnlineRetail$TransactionValue <- OnlineRetail$Quantity * OnlineRetail$UnitPrice

# Aggregate the total revenue generated by each product
product_revenue <- aggregate(TransactionValue ~ StockCode, data = OnlineRetail, FUN = sum)
highest_revenue_product <- product_revenue[which.max(product_revenue$TransactionValue), "StockCode"]
print(paste("Product with the Highest Revenue:", highest_revenue_product))
```

```
## [1] "Product with the Highest Revenue: DOT"
```

## Question 12 : How many unique customers are represented in the dataset? You can use unique() and length() functions.

```
# Assuming you have the 'Online_Retail' dataset with a 'CustomerID' column

# Extract unique customer IDs
unique_customers <- unique(OnlineRetail$CustomerID)

# Calculate the number of unique customers
num_unique_customers <- length(unique_customers)

# Display the result
print(paste("Number of Unique Customers:", num_unique_customers))
```

```
## [1] "Number of Unique Customers: 4373"
```